

BLOG SUBMISSION

HR Analytics Project - Understanding the Attrition in HR



Submitted By:

M Ashok Kumar

ACKNOWLEDGMENT

In this blog, I will be analysing the HR Analytics Project- Understanding the Attrition in HR using Machine Learning dataset, using essential exploratory data analysis techniques and also, I will be performing some data visualizations to better understand our data.

In the dataset, there are many columns like Age, Attrition, Daily Rate, Department, Education and so on. By doing data pre-processing, data analysis, feature selection, and many other techniques we built our cool and fancy machine learning model. And at the end, we applied many ML algorithms to get the very good accuracy of our model.

Many thanks to Data Trained for providing me required knowledge with this project to understand the Realtime Field work present in the Data Science Industry.

I am very Thankful to the entire team and also my family members for encouraging and supporting me in submitting this blog.

ABSTRACT

Every year a lot of companies hire a number of employees. The companies invest time and money in training those employees, not just this but there are training programs within the companies for their existing employees as well. The aim of these programs is to increase the effectiveness of their employees. But where HR Analytics fit in this? and is it just about improving the performance of employees?

TAKEAWAYS FROM THE BLOG

In this article, we do prediction using machine learning which leads to the below takeaways:

1. EDA: Learn the complete process of EDA
2. Data analysis: Learn to withdraw some insights from the dataset both mathematically and visualize it.
3. Data visualization: Visualizing the data to get better insight from it.
4. Feature engineering: We will also see what kind of stuff we can do in the feature engineering part.

PROBLEM STATEMENT:

Every year a lot of companies hire a number of employees. The companies invest time and money in training those employees, not just this but there are training programs within the companies for their existing employees as well. The aim of these programs is to increase the effectiveness of their employees. But where HR Analytics fit in this? and is it just about improving the performance of employees?

HR Analytics:

Human resource analytics (HR analytics) is an area in the field of analytics that refers to applying analytic processes to the human resource department of an organization in the hope of improving employee performance and therefore getting a better return on investment. HR analytics does not just deal with gathering data on employee efficiency. Instead, it aims to provide insight into each process by gathering data and then using it to make relevant decisions about how to improve these processes.

Attrition in HR:

Attrition in human resources refers to the gradual loss of employee's overtime. In general, relatively high attrition is problematic for companies. HR professionals often assume a leadership role in designing company compensation programs, work culture, and motivation systems that help the organization retain top employees.

How does Attrition affect companies? and how does HR Analytics help in analysing attrition? We will discuss the first question here and for the second question, we will write the code and try to understand the process step by step.

Attrition affecting Companies

A major problem in high employee attrition is its cost to an organization. Job postings, hiring processes, paperwork, and new hire training are some of the common expenses of losing employees and replacing them. Additionally, regular employee turnover prohibits your organization from increasing its collective knowledge base and experience over time. This is especially concerning if your business is customer-facing, as customers often prefer to interact with familiar people. Errors and issues are more likely if you constantly have new workers.

Size of the Dataset:

The Dataset Contains 1470 rows × 35 columns.

ABOUT THE DATASET

You can find the dataset in the link below:

- https://github.com/dsrscientist/IBM_HR_Attrition_Rate_Analytics

Importing Important Libraries:

We need some libraries to be imported to work upon the dataset, we would import the dataset by using pandas read_csv method.

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import warnings
warnings.filterwarnings("ignore")
from sklearn.model_selection import train_test_split, GridSearchCV, RandomizedSearchCV, cross_val_score
from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestClassifier, AdaBoostClassifier
from sklearn.tree import DecisionTreeClassifier
from sklearn.naive_bayes import GaussianNB
from sklearn.svm import SVC
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import classification_report, roc_auc_score, confusion_matrix, accuracy_score
```

Loading Data Set into a Desired variable:

Here I am loading the Dataset into the variable hr_df.

```
hr_df = pd.read_csv("Hr Analysis.csv")
hr_df
```

Here we have imported the dataset into our callable variable and the dataset contains 1470 columns and 35 rows.

Exploratory Data Analysis:

Before you start a machine learning project, it's important to ensure that the data is ready for modelling work.

Exploratory Data Analysis (EDA) ensures the readiness of the data for Machine Learning.

In fact, EDA is primarily used to see what data can reveal beyond the formal modelling or hypothesis testing task and provides a better understanding of data set variables and the relationships between them.

Let's do the EDA using the statistical techniques.

Getting detailed info about the Dataset:

```
hr_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1470 entries, 0 to 1469
Data columns (total 35 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Age                                   1470 non-null   int64
1   Attrition                           1470 non-null   object
2   BusinessTravel                      1470 non-null   object
3   DailyRate                           1470 non-null   int64
4   Department                          1470 non-null   object
5   DistanceFromHome                   1470 non-null   int64
6   Education                           1470 non-null   int64
7   EducationField                      1470 non-null   object
8   EmployeeCount                      1470 non-null   int64
9   EmployeeNumber                     1470 non-null   int64
10  EnvironmentSatisfaction             1470 non-null   int64
11  Gender                             1470 non-null   object
12  HourlyRate                          1470 non-null   int64
13  JobInvolvement                     1470 non-null   int64
14  JobLevel                           1470 non-null   int64
15  JobRole                            1470 non-null   object
16  JobSatisfaction                    1470 non-null   int64
17  MaritalStatus                      1470 non-null   object
18  MonthlyIncome                     1470 non-null   int64
19  MonthlyRate                        1470 non-null   int64
20  NumCompaniesWorked                 1470 non-null   int64
21  Over18                            1470 non-null   object
22  OverTime                           1470 non-null   object
23  PercentSalaryHike                  1470 non-null   int64
24  PerformanceRating                  1470 non-null   int64
25  RelationshipSatisfaction            1470 non-null   int64
26  StandardHours                     1470 non-null   int64
27  StockOptionLevel                   1470 non-null   int64
28  TotalWorkingYears                  1470 non-null   int64
29  TrainingTimesLastYear              1470 non-null   int64
30  WorkLifeBalance                    1470 non-null   int64
31  YearsAtCompany                     1470 non-null   int64
32  YearsInCurrentRole                 1470 non-null   int64
33  YearsSinceLastPromotion            1470 non-null   int64
34  YearsWithCurrManager               1470 non-null   int64
dtypes: int64(26), object(9)
memory usage: 402.1+ KB
```

There are 26 integer columns and 9 object type columns are present in the dataset

and there are no null values present in the dataset.

Let's Check the statistical Data for the Dataset:

Here, we use `describe()` to get the statistical values of the Dataset.

```
hr_df.describe()
```

	Age	DailyRate	DistanceFromHome	Education	EmployeeCount	EmployeeNumber	EnvironmentSatisfaction	HourlyRate	JobInvolvement
count	1470.000000	1470.000000	1470.000000	1470.000000	1470.0	1470.000000	1470.000000	1470.000000	1470.000000
mean	36.923810	802.485714	9.192517	2.912925	1.0	1024.865306	2.721769	65.891156	2.729932
std	9.135373	403.509100	8.106864	1.024165	0.0	602.024335	1.093082	20.329428	0.711561
min	18.000000	102.000000	1.000000	1.000000	1.0	1.000000	1.000000	30.000000	1.000000
25%	30.000000	465.000000	2.000000	2.000000	1.0	491.250000	2.000000	48.000000	2.000000
50%	36.000000	802.000000	7.000000	3.000000	1.0	1020.500000	3.000000	66.000000	3.000000
75%	43.000000	1157.000000	14.000000	4.000000	1.0	1555.750000	4.000000	83.750000	3.000000
max	60.000000	1499.000000	29.000000	5.000000	1.0	2068.000000	4.000000	100.000000	4.000000

8 rows × 10 columns

By observing the data:


There is a huge difference for 75th percentile and maximum values for EmployeeNumber. so there may be outliers present in the dataset.

Let's Analyse the single variables separately by using the Univariate Analysis

Checking the categorical columns and Plotting graphs for better insight of Data Distribution

```
#Checking the target column  
sns.countplot(x=hr_df['Attrition'])  
print(hr_df['Attrition'].value_counts())
```

```
No      1233  
Yes      237  
Name: Attrition, dtype: int64
```




Here we can see that out of 1470 employees, 237 employees which is around 16% of total employees left their job due to some reasons.

whereas other 1233 employees, which is 84% of the employees preferred to continue in their job at the company.

Also, here we can observe that there is huge difference between two types of Attrition. So, the data is imbalanced. So, we will apply SMOTE analysis before ML of final model.


```
sns.countplot(x=hr_df['BusinessTravel'])  
print(hr_df['BusinessTravel'].value_counts())
```

```
Travel_Rarely      1043  
Travel_Frequently   277  
Non-Travel         150  
Name: BusinessTravel, dtype: int64
```




Here, business travel is divided in 3 categories with least number of non-travellers.

It has maximum number of employees who travels rarely.

```
sns.countplot(x=hr_df['Department'])  
print(hr_df['Department'].value_counts())
```

```
Research & Development   961  
Sales                    446  
Human Resources          63  
Name: Department, dtype: int64
```




Out of 3 different departments in the dataset Research & Development has maximum number of employees.

```
sns.countplot(x=hr_df['EducationField'])  
print(hr_df['EducationField'].value_counts())
```

Life Sciences	606
Medical	464
Marketing	159
Technical Degree	132
Other	82
Human Resources	27


Name: EducationField, dtype: int64



Most of the employees are from life sciences background and least are from human resource.

```
sns.countplot(x=hr_df['Gender'])
print(hr_df['Gender'].value_counts())
```


```
Male      882
Female    588
Name: Gender, dtype: int64
```



We can Observe that male employee has more attrition than female employees.

```
sns.countplot(x=hr_df['JobRole'])
print(hr_df['JobRole'].value_counts())
```


```
Sales Executive      326
Research Scientist   292
Laboratory Technician 259
Manufacturing Director 145
Healthcare Representative 131
Manager              102
Sales Representative   83
Research Director     80
Human Resources       52
Name: JobRole, dtype: int64
```



Most of the employee are from Sales department. And least are from HR.

```
sns.countplot(x=hr_df['MaritalStatus'])  
print(hr_df['MaritalStatus'].value_counts())
```


```
Married    673  
Single     470  
Divorced   327  
Name: MaritalStatus, dtype: int64
```



Most of the employees are married.

```
sns.countplot(x=hr_df['OverTime'])  
print(hr_df['OverTime'].value_counts())
```

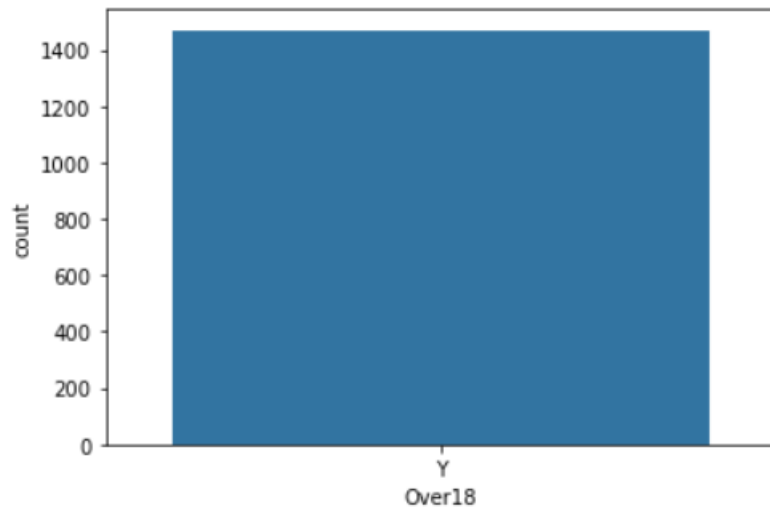
```
No      1054  
Yes      416  
Name: OverTime, dtype: int64
```



Only 416 employees work overtime.

```
sns.countplot(x=hr_df['Over18'])
```

```
<AxesSubplot:xlabel='Over18', ylabel='count'>
```



All the employees are above 18.

so, we can drop this column as there is no use of this.

Analysis of the Rating Features JobSatisfaction

EnvironmentSatisfaction

RelationshipSatisfaction

JobInvolvement

WorkLifeBalance

PerformanceRating

```

fig = plt.figure()

ax1 = fig.add_subplot(221)
ax2 = fig.add_subplot(222)
ax3 = fig.add_subplot(223)
ax4 = fig.add_subplot(224)

labels = 'Low', 'Medium', 'High', 'Very High'

hr_df['JobSatisfaction'].astype(str).value_counts().plot(kind='pie',
    figsize=(15, 6),
    autopct='%1.1f%%',
    startangle=90,
    shadow=True,
    labels=None, ax=ax1) # add to subplot 2
ax1.set_title('Rating of Job Satisfaction by Employees')
fig.legend(labels=labels, loc='center')

hr_df['EnvironmentSatisfaction'].astype(str).value_counts().plot(kind='pie',
    figsize=(15, 6),
    autopct='%1.1f%%',
    startangle=90,
    shadow=True,
    labels=None, ax=ax2)
ax2.set_title('Rating of Environmental Satisfaction by Employees')

hr_df['RelationshipSatisfaction'].astype(str).value_counts().plot(kind='pie',
    figsize=(15, 6),
    autopct='%1.1f%%',
    startangle=90,
    shadow=True,
    labels=None, ax=ax3)
ax3.set_title('Rating of Relationship Satisfaction by Employees')

hr_df['JobInvolvement'].astype(str).value_counts().plot(kind='pie',
    figsize=(15, 6),
    autopct='%1.1f%%',
    startangle=90,
    shadow=True,
    labels=None, ax=ax4)
ax4.set_title('Rating of Job Involvement by Employees')

plt.show()


```



Here we can see the rating features are low for each cases which also effects the performance of the employees.

Not Getting involved in their job


```
plt.figure(figsize=(10,5))
sns.countplot(x='Age',hue='Attrition',data=hr_df)
plt.title('Agewise Attrition Rate')
print(hr_df.groupby('Age')['Attrition'].value_counts())
```



Let's Check which gender has more attrition.

```
plt.figure(figsize=(10,5))
sns.countplot(x='Gender',hue='Attrition',data=hr_df)
print(hr_df.groupby('Gender')['Attrition'].value_counts())
```

```
Gender  Attrition
Female  No         501
        Yes         87
Male    No         732
        Yes        150
Name: Attrition, dtype: int64
```



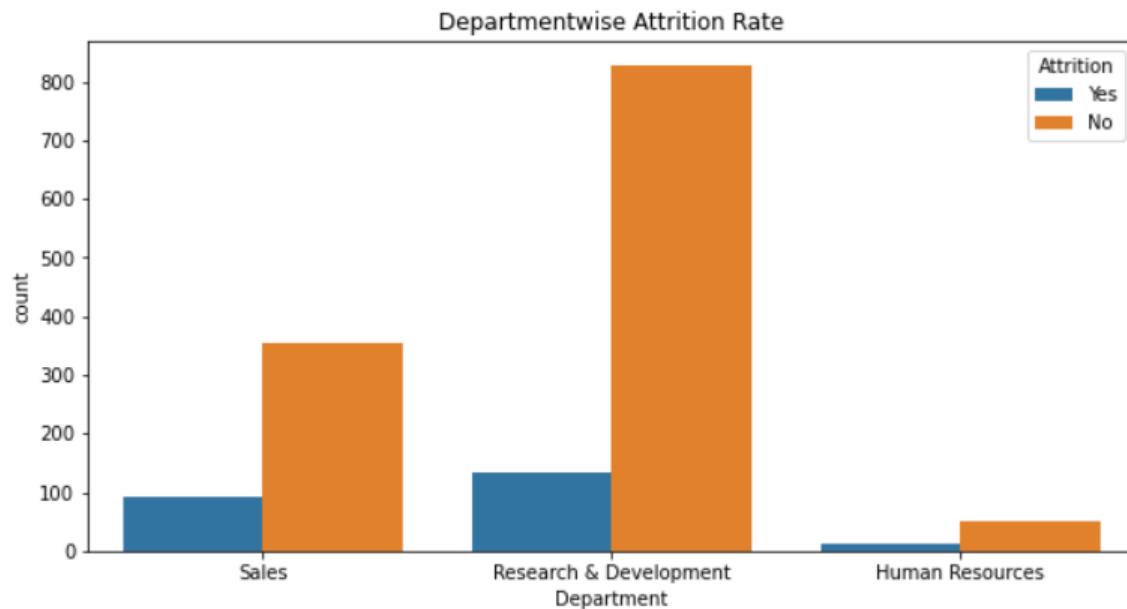
Here, we can see attrition does not much depend on sex of the employees.

Let's check which department has highest attrition.

```
plt.figure(figsize=(10,5))
sns.countplot(x='Department',hue='Attrition',data=hr_df)
plt.title('Departmentwise Attrition Rate')
print(hr_df.groupby('Department')['Attrition'].value_counts())
```

Department	Attrition	
Human Resources	No	51
	Yes	12
Research & Development	No	828
	Yes	133
Sales	No	354
	Yes	92

Name: Attrition, dtype: int64

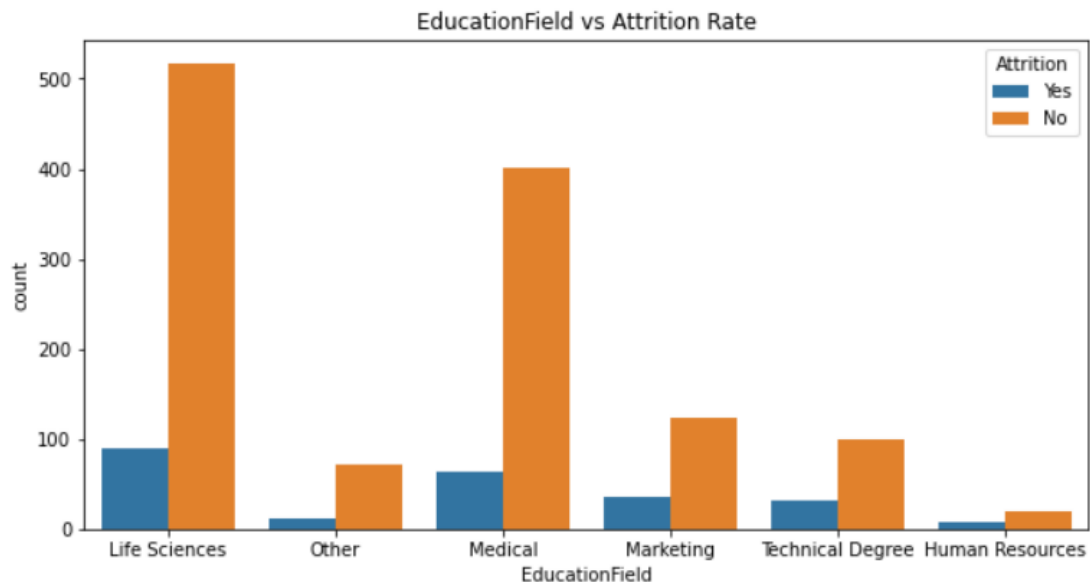


From the above graph we can see that Highest attrition is in Research & Development department.

```
plt.figure(figsize=(10,5))
sns.countplot(x='EducationField',hue='Attrition',data=hr_df)
plt.title('EducationField vs Attrition Rate')
print(hr_df.groupby('EducationField')['Attrition'].value_counts())
```

EducationField	Attrition	
Human Resources	No	20
	Yes	7
Life Sciences	No	517
	Yes	89
Marketing	No	124
	Yes	35
Medical	No	401
	Yes	63
Other	No	71
	Yes	11
Technical Degree	No	100
	Yes	32

Name: Attrition, dtype: int64



We can see that the people who are in life sciences have more attrition rate.


```

yac = hr_df.groupby("YearsAtCompany")['Attrition'].value_counts(normalize=False).unstack()

yac.plot(kind='bar', stacked=False, figsize=(10,5))

plt.title('Years At Company of Employee')
plt.ylabel('Number of Employees')
plt.show()

```



From the above graph It is observed that the newly arriving employees quit their jobs most.

```

plt.figure(figsize=(20,5))
sns.countplot(x='MonthlyIncome', hue='Attrition', data=hr_df)
plt.title('MonthlyIncome vs Attrition Rate')
print(hr_df.groupby('MonthlyIncome')['Attrition'].value_counts())

```

```

MonthlyIncome  Attrition
1009           Yes      1
1051           No       1
1052           No       1
1081           Yes      1
1091           Yes      1
..
19859          Yes      1
19926          No       1
19943          No       1
19973          No       1
19999          No       1
Name: Attrition, Length: 1388, dtype: int64

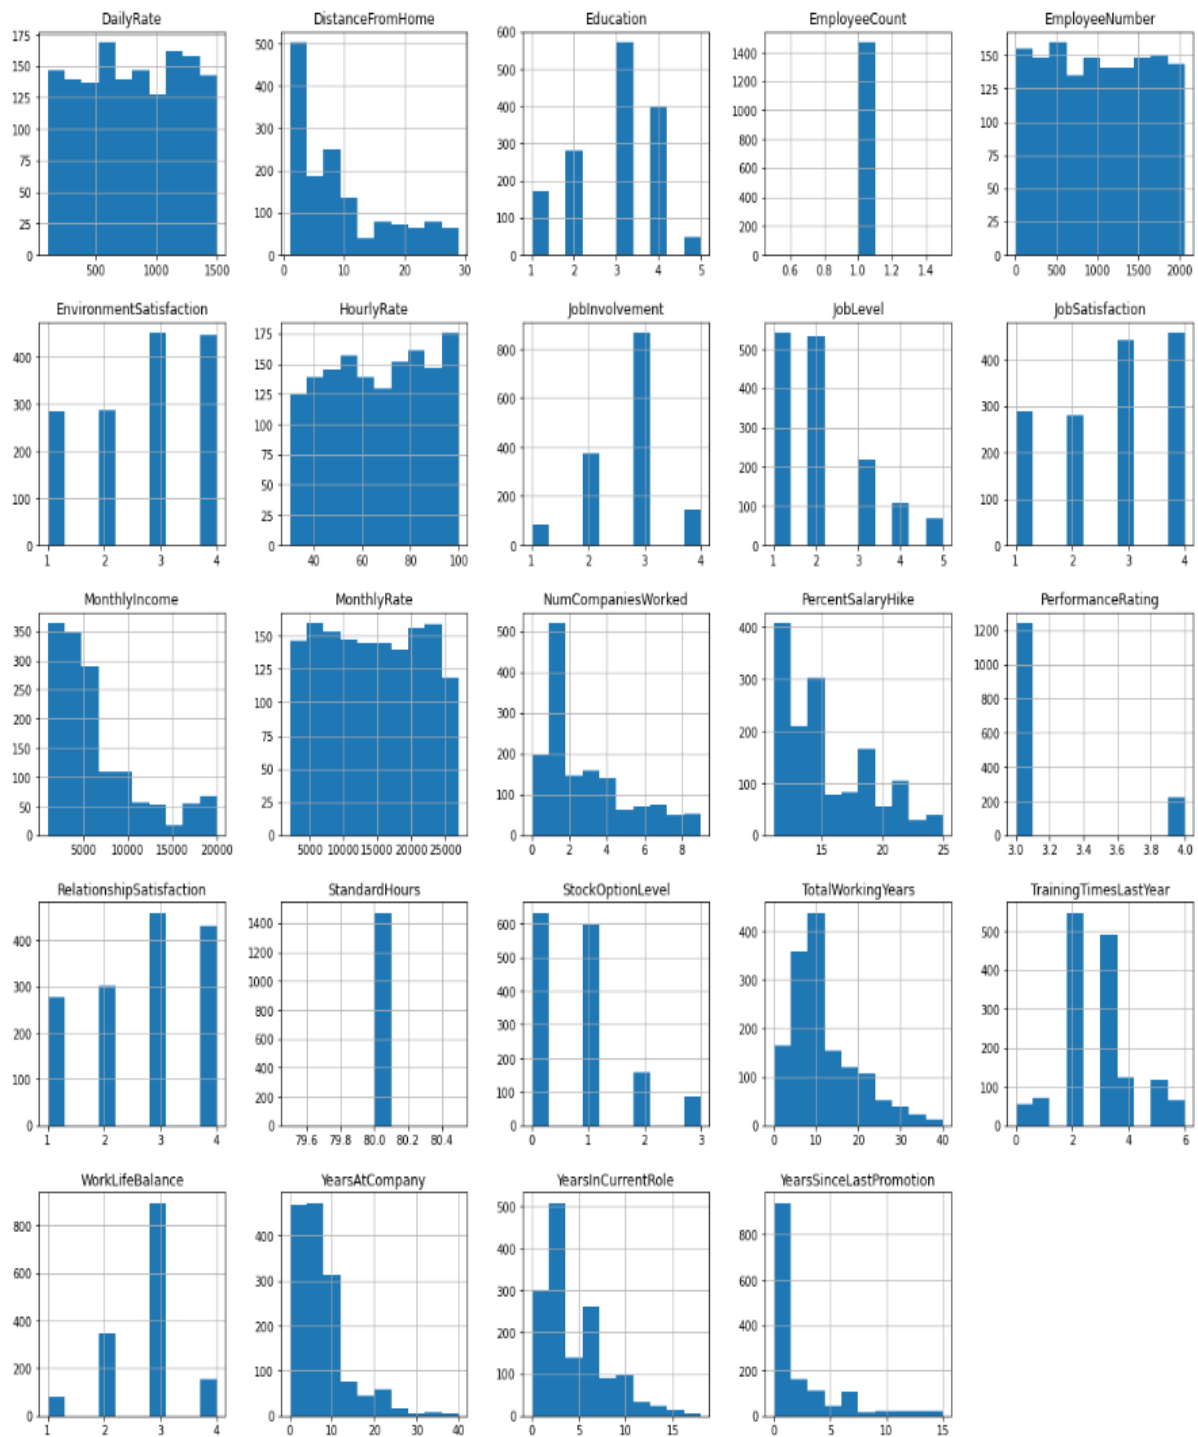
```



Employees who left their jobs tend to have low monthly income than those who continued their job in the company.

Let's check the data with Multivariate analysis:

```
# distribution of the data  
hr_df.iloc[:,1:-1].hist(figsize=(20,20));
```



Converting categorical values to numerical values using label encoder:

```
# use label encoder to change data type in type and region columns
from sklearn.preprocessing import LabelEncoder
le=LabelEncoder()
list=['Attrition','BusinessTravel','Department','EducationField','Gender','JobRole','MaritalStatus','OverTime']
for val in list:
    hr_df[val]=le.fit_transform(hr_df[val].astype(str))
```

For Attrition column, "Yes":1,"No":0

For BusinessTravel column, "non-Travel":0,"Travel_Frequently":1,"Travel_Rarely":2

For Department column, "Human Resources":0,"Research & Development":1,"Sales":2.

For Gender column,"Female":0,"Male":1

For MaritalStatus column, "Divorced":0,"Married":1,"Single":2

For OverTime column,"Yes":0,"No":1

For JobRole column,'Sales Executive': 7,'Research Scientist': 5,laboratory Technician': 2,'Manufacturing Director': 3, 'Healthcare Representative': 1,'Manager': 4,'Sales Representative': 8,'Research Director': 6,'Human Resources': 0

For EducationField column,Life Sciences': 2,'Other': 1,'Medical':4 ,'Marketing': 3,'Technical Degree': 5, 'HumanResources': 0

Over18 and EmployeeCount have only one value thus it will not provide any information about the data. Considering EmployeeNumber is emp ID thus deleting it.

```
hr_df.drop(columns=["Over18","EmployeeCount","EmployeeNumber","StandardHours"],axis=1, inplace=True)
```


```
hr_df.head(2)
```

	Age	Attrition	BusinessTravel	DailyRate	Department	DistanceFromHome	Education	EducationField	EnvironmentSatisfaction	Gender	...	PerformanceRating
0	41	1	2	1102	2	1	2	1	2	0	...	3
1	49	0	1	279	1	8	1	1	3	1	...	4

2 rows × 31 columns

Now there are 31 columns left.

Now, Let's check the correlation:



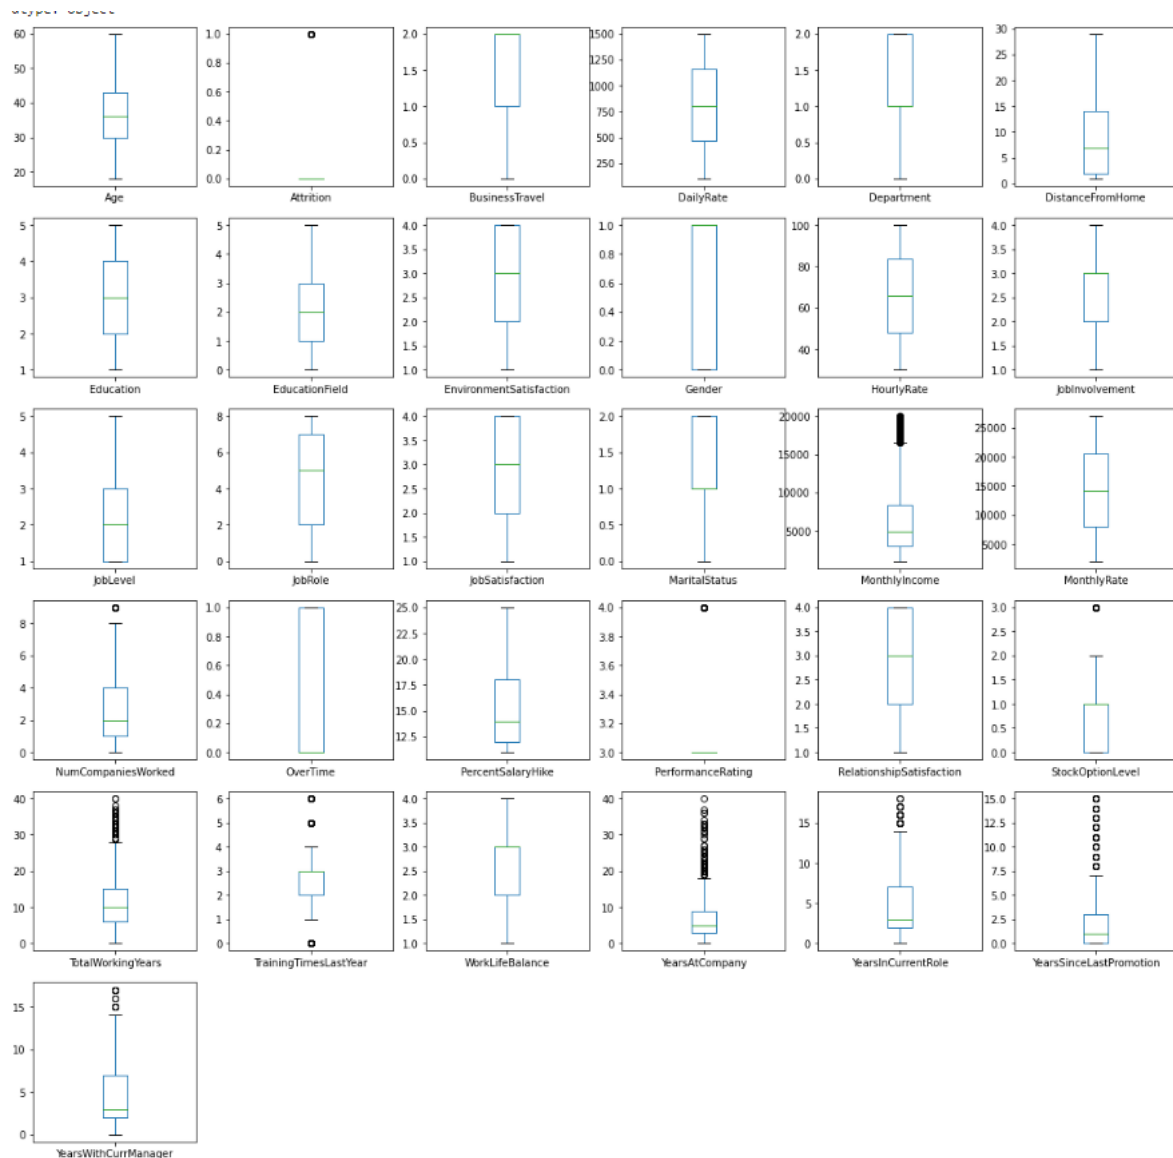
None of the column is highly correlated with Attrition. So many columns are negatively correlated with Attrition.

Job role is highly correlated with department, job level is highly positively correlated with monthlyincome and totalworkingyear, total working year is correlated with Age, job level, monthly income.

year of company correlated with year at current role and year with current manager.

Let's check whether there are outliers present in the data or not.

```
hr_df.plot(kind='box',subplots=True,layout=(6,6),figsize=(20,20))
```



There are outliers in MonthlyIncome, Stock option Level, Total Working years, Training Times last year, Years at company, years in current role and Years with current manager.

Let's Remove the outliers

```
from scipy.stats import zscore

z=np.abs(zscore(hr_df))
z
```

Let's check the shape of the data.

```
z.shape
```

```
(1470, 31)
```

```
threshold=3
print(np.where(z>3))
```

```
(array([ 28,  45,  62,  62,  63,  64,  85,  98,  98, 110, 123,
        123, 123, 126, 126, 126, 153, 178, 187, 187, 190, 190,
        218, 231, 231, 237, 237, 270, 270, 281, 326, 386, 386,
        401, 411, 425, 425, 427, 445, 466, 473, 477, 535, 561,
        561, 584, 592, 595, 595, 595, 616, 624, 635, 653, 653,
        677, 686, 701, 716, 746, 749, 752, 799, 838, 861, 861,
        875, 875, 894, 914, 914, 918, 922, 926, 926, 937, 956,
        962, 976, 976, 1008, 1024, 1043, 1078, 1078, 1086, 1086, 1093,
        1111, 1116, 1116, 1135, 1138, 1138, 1156, 1184, 1221, 1223, 1242,
        1295, 1301, 1301, 1303, 1327, 1331, 1348, 1351, 1401, 1414, 1430],
      dtype=int64), array([30, 29, 27, 29, 28, 29, 24, 24, 27, 29, 28, 29, 30, 24, 27, 29, 30, 24,
        29, 24, 30, 27, 28, 29, 28, 30, 27, 29, 24, 27, 28, 29, 29, 30, 24,
        27, 27, 29, 29, 24, 28, 27, 27, 29, 27, 30, 29, 27, 24, 27, 29, 30,
        24, 30, 27, 29, 27, 30, 29, 28, 28, 27, 29, 29, 29, 27, 29, 29, 30,
        24, 27, 29, 27, 29, 29, 30, 29, 24, 27, 28, 29, 29, 28, 24, 29, 30,
        27, 29, 29, 27, 24, 27, 27, 27, 29, 29, 24, 29, 29, 29, 29, 24, 29,
        29, 28, 29, 30, 28, 24, 29, 28], dtype=int64))
```

```
print(hr_df.shape)
print(hr_new.shape)
```

```
(1470, 31)
```

```
(1387, 31)
```

From the above we can understand that total dataset contains 1470 rows and 31 columns.

But after removing the outliers, the data contains 1387 rows and 31 columns.

So, there is a loss of 5.71% of the data.


Let's Handle the imbalance data by using SMOTE Analysis.

```
# Separating data into X and y
x_new = hr_new.drop("Attrition",axis =1)
y_new= hr_new["Attrition"]
```

```
# implementing oversampling for handling imbalance data
from imblearn.over_sampling import SMOTE
smt=SMOTE()
sm_x,sm_y=smt.fit_resample(x_new,y_new)
#y.value_counts()
sm_y.value_counts()
```

```
0    1158
1    1158
Name: Attrition, dtype: int64
```

```
plt.figure(figsize=(5,5))
sns.countplot(sm_y)
plt.show()
```



Now the data is Balanced.

Let's do the Modelling:

```
maxAcc = 0
maxRS = 0

for i in range(1,200):
    x_train,x_test,y_train,y_test=train_test_split(sm_x,sm_y,random_state=i,test_size=0.20)
    LR = LogisticRegression()
    LR.fit(x_train,y_train)
    predrf= LR.predict(x_test)
    acc=accuracy_score(y_test,predrf)
    if acc>maxAcc:
        maxAcc=acc
        maxRS=i
print('Best score is',maxAcc,'on Random State',maxRS)
```

Best score is 0.8685344827586207 on Random State 75

Let's Make input and output variables into train and test data.

```
from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test=train_test_split(sm_x,sm_y,random_state=75,test_size=0.20)
```

Let's check the size of the test and train:

```
print(y_train.shape,'\t',y_test.shape)
```

(1852,) (464,)

Let's check with different algorithms:

```
model=[LogisticRegression(),GaussianNB(),SVC(),DecisionTreeClassifier(),KNeighborsClassifier(),]

for m in model:
    m.fit(x_train,y_train)
    m.score(x_train,y_train)
    predm=m.predict(x_test)
    print('Accuracy score of',m, 'is:')
    print(accuracy_score(y_test,predm))
    print(confusion_matrix(y_test,predm))
    print(classification_report(y_test,predm))
    print('\n')
```

Accuracy score of LogisticRegression() is:

0.8685344827586287

[[212 31]

[30 191]]

	precision	recall	f1-score	support
0	0.88	0.87	0.87	243
1	0.86	0.86	0.86	221
accuracy			0.87	464
macro avg	0.87	0.87	0.87	464
weighted avg	0.87	0.87	0.87	464

Accuracy score of GaussianNB() is:

0.75

[[169 74]

[42 179]]

	precision	recall	f1-score	support
0	0.80	0.70	0.74	243
1	0.71	0.81	0.76	221
accuracy			0.75	464
macro avg	0.75	0.75	0.75	464
weighted avg	0.76	0.75	0.75	464

Accuracy score of SVC() is:

0.9267241379310345

[[224 19]

[15 206]]

	precision	recall	f1-score	support
0	0.94	0.92	0.93	243
1	0.92	0.93	0.92	221
accuracy			0.93	464
macro avg	0.93	0.93	0.93	464
weighted avg	0.93	0.93	0.93	464

Accuracy score of DecisionTreeClassifier() is:

0.8405172413793104

[[196 47]

[27 194]]

	precision	recall	f1-score	support
0	0.88	0.81	0.84	243
1	0.80	0.88	0.84	221
accuracy			0.84	464
macro avg	0.84	0.84	0.84	464
weighted avg	0.84	0.84	0.84	464

Accuracy score of KNeighborsClassifier() is:

0.8577586206896551

[[178 65]

[1 220]]

	precision	recall	f1-score	support
0	0.99	0.73	0.84	243
1	0.77	1.00	0.87	221
accuracy			0.86	464
macro avg	0.88	0.86	0.86	464
weighted avg	0.89	0.86	0.86	464

We got the highest accuracy with SVC.

```
# Finding out best paramter using GridsearchCV
from sklearn.model_selection import GridSearchCV
parameters={'kernel':['linear','rbf','poly','sigmoid'],'C':[0,1,2,3,4,5,6,7,8,9,10,11,12]}
svc=SVC()
gs=GridSearchCV(svc,parameters)
gs.fit(sm_x,sm_y)
```

```
GridSearchCV(estimator=SVC(),
              param_grid={'C': [0, 1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12],
                           'kernel': ['linear', 'rbf', 'poly', 'sigmoid']})
```

```
print(gs.best_params_)
```

```
{'C': 10, 'kernel': 'rbf'}
```

```
#Using SVC model with best results
sv=SVC(kernel='rbf',C=10)
sv.fit(x_train,y_train)
sv.score(x_train,y_train)
predsv=sv.predict(x_test)
print('Accuracy score of',sv, 'is:')
print(accuracy_score(y_test,predsv))
print(confusion_matrix(y_test,predsv))
print(classification_report(y_test,predsv))
print('\n')
```

Accuracy score of SVC(C=10) is:

0.9482758620689655

[[225 18]

[6 215]]

	precision	recall	f1-score	support
0	0.97	0.93	0.95	243
1	0.92	0.97	0.95	221
accuracy			0.95	464
macro avg	0.95	0.95	0.95	464
weighted avg	0.95	0.95	0.95	464

Let's use Ensemble Technique to boost up score:

```
from sklearn.ensemble import RandomForestClassifier

rf=RandomForestClassifier(n_estimators=100,random_state=75,criterion='gini')
#RandomForestClassifier(100)---Default
rf.fit(x_train,y_train)
predrf=rf.predict(x_test)
print(accuracy_score(y_test,predrf))
print(confusion_matrix(y_test,predrf))
print(classification_report(y_test,predrf))
```

0.9461206896551724

[[230 13]

[12 209]]

	precision	recall	f1-score	support
0	0.95	0.95	0.95	243
1	0.94	0.95	0.94	221
accuracy			0.95	464
macro avg	0.95	0.95	0.95	464
weighted avg	0.95	0.95	0.95	464

```
from sklearn.ensemble import AdaBoostClassifier
ad=AdaBoostClassifier(n_estimators=100,random_state=75,base_estimator=sv,algorithm='SAMME',learning_rate=0.01)
ad.fit(x_train,y_train)
ad_pred=ad.predict(x_test)
print(accuracy_score(y_test,ad_pred))
print(confusion_matrix(y_test,ad_pred))
print(classification_report(y_test,ad_pred))
```

0.47629310344827586

[[0 243]

[0 221]]

	precision	recall	f1-score	support
0	0.00	0.00	0.00	243
1	0.48	1.00	0.65	221
accuracy			0.48	464
macro avg	0.24	0.50	0.32	464
weighted avg	0.23	0.48	0.31	464

```
# Finding out best paramter using GridsearchCV
from sklearn.model_selection import GridSearchCV
parameters={'random_state':range(35,100)}
gc=GridSearchCV(sv,parameters)
gc.fit(sm_x,sm_y)
gc.best_params_
```

```
{'random_state': 35}
```

Hyperparameter Tuning

```
RF = RandomForestClassifier() # Randomforest instance
from sklearn.model_selection import RandomizedSearchCV,cross_val_score

# we are checking best parameter which suited most for the model

para = {
    "criterion":["gini","entropy"],
    "max_features" : ["auto", "sqrt", "log2"],
    "min_samples_split" :[2,4,6,8],
    "max_depth" :[2,3,4,5],
    "min_samples_leaf":[3,5,8,10],
    "max_samples" :[0.2,0.3,0.4,0.5]
}

grid = RandomizedSearchCV(RF, para, scoring= "accuracy" )

grid.fit(x_train,y_train)
print("score",grid.best_score_)
print("best parameter",grid.best_params_)
print("best estimator",grid.best_estimator_)
grid_pred=grid.best_estimator_.predict(x_test) # predicting with best parameters
```

score 0.8326174692212428

best parameter {'min_samples_split': 4, 'min_samples_leaf': 3, 'max_samples': 0.4, 'max_features': 'log2', 'max_depth': 5, 'criterion': 'gini'}
 best estimator RandomForestClassifier(max_depth=5, max_features='log2', max_samples=0.4,
 min_samples_leaf=3, min_samples_split=4)

Cross Validation

```
best_parameter_RF = RandomForestClassifier(min_samples_split= 4, min_samples_leaf = 3, max_samples =0.4, max_features = 'log2', max_depth = 5, crite

for i in range(2,7):
    cv = cross_val_score(best_parameter_RF,sm_x,sm_y,cv=i)
    print(f'at CV {i} The mean is {cv.mean()} and the SD is {cv.std()}')
```

at CV 2 The mean is 0.8385146804835923 and the SD is 0.03108808290155446
 at CV 3 The mean is 0.8337651122625216 and the SD is 0.04518644370276731
 at CV 4 The mean is 0.8354922279792746 and the SD is 0.05899275910295336
 at CV 5 The mean is 0.8247598123184628 and the SD is 0.07208950384366851
 at CV 6 The mean is 0.8380829015544041 and the SD is 0.06157506645991527

Observation

since Randomforest worked well out of all other model, so we have done the hyperparameter tuning to set the best parameter for final model. Now i have checked the best CV as well that at level of CV is generated the best score and we have found CV 6 is at best


```
# Cross validate of RandomForestClassifier using cv=6
from sklearn.model_selection import cross_val_score
score=cross_val_score(rf,sm_x,sm_y,cv=6,scoring='accuracy')
print('Score:', score)
print('Mean Score:', score.mean())
print('Standard Deviation:', score.std())
```

Score: [0.73834197 0.92746114 0.94041451 0.9507772 0.94041451 0.93005181]
 Mean Score: 0.9045768566493955
 Standard Deviation: 0.07473268822844518

```
auc_score=roc_auc_score(y_test,rf.predict(x_test))
print(auc_score)
```

0.9461017075396161

```
plt.plot([0,1],[0,1], 'k--')
plt.plot(fpr, tpr, label='RandomForestRegressor')
plt.xlabel('False positive rate')
plt.ylabel('True positive rate')
plt.title('RandomForestRegressor')
plt.show()
```



I found out best result with RandomForestRegressor by using various techniques. So I will save RandomForestRegressor as my final model.

Conclusion:

So, as we saw that we have done a complete EDA process, getting data insights, feature engineering, and data visualization as well so after all these steps one can go for the prediction using machine learning model-making steps.

We have training and test file separately available with us. All the independent variables are categorical in nature and the dependent variable of train data i.e., the attrition is the classical data type. So, I applied the regression method for prediction.

I found out best result with RandomForestRegressor by using various techniques. So I will save RandomForestRegressor as my final model.

I hope this article helped you to understand Data Analysis, Data Preparation, and Model building approaches in a much simpler way.

Thank you for reading this blog